

# Data Analysis with Python Pandas Training v6



Trainer: Marcus Lee



Website: [www.tertiarycourses.com.sg](http://www.tertiarycourses.com.sg)  
Email: enquiry@tertiaryinfotech.com<sup>1</sup>

# About the Trainer

Marcus Lee has his degree in Computer Science and a background in Statistics from the University of Otago. Before returning to Singapore, he analysed vacation data provided by the New Zealand Board of Tourism to determine the favourite activities of Australian, Japanese, and German tourists in New Zealand. In addition to a vast number of other demographics statistics, he has been able to provide significant advice to the board on how to promote tourism in New Zealand. His core specialization skills are Java, R, Statistical Analysis, Machine Learning, NumPy, Scikit, and Network Management. He has also a fair amount of experience in C, C++, and Python



# Ground Rules

- Set your mobile phone to silent mode
- Participate actively in the class. No question is stupid.
- Mutual respect. Agree to disagree.
- One conversation at one time.
- Be punctual. Back from breaks on time.
- Exit the class silently if you need to step out for phone call, toilet break etc.
- 75% attendance is required

# Ground Rules for Virtual Training

- Upon entering, mute your mic and turn on the video. Use a headset if you can
- Use the 'raise hand' function to indicate when you want to speak
- Participant actively. Feel free to ask questions on the chat whenever.
- Facilitators can use breakout rooms for private sessions.



# Guidelines for Facilitators

1. Once all the participants are in and introduce themselves
2. Goto gallery mode, take a snapshot of the class photo - makes sure capture the date and time
3. Start the video recording (only for WSQ courses)
4. Continue the class
5. Before the class end on that day, take another snapshot of the class photo - makes sure capture the date and time
6. For NRIC verification, facilitator to create breakout room for individual participant to check (only for WSQ courses)
7. Before the assessment start, take another snapshot of the class photo - makes sure capture the date and time (only for WSQ courses)
8. For Oral Questioning assessment, facilitator to create breakout room for individual participant to OQ (only for WSQ courses)
9. End the video recording and upload to cloud (only for WSQ courses)
10. Assessor to send all the assessment records, assessment plan and photo and video to the staff (only for WSQ courses).

# Prerequisite

This course assumes the following knowledge

- Basic Python

# Agenda

## Topic 1 Data Preparation

- Data Analytics with Pandas
- Pandas DataFrame and Series
- Import and Export Data
- Filter and Slice Data
- Clean Data

## Topic 2 Data Transformation

- Join Data
- Transform Data
- Aggregate Data

## Topic 3 Data Visualization

- Data Visualization with Matplotlib and Seaborn
- Visualize Statistical Relationships with Scatter Plot
- Visualize Categorical Data with Bar Plot
- Visualize Correlation with Pair Plot and Heatmap
- Visualize Linear Relationships with Regression

# Agenda

## Topic 4 Data Analysis

- Statistical Data Analysis
- Time Series Analysis



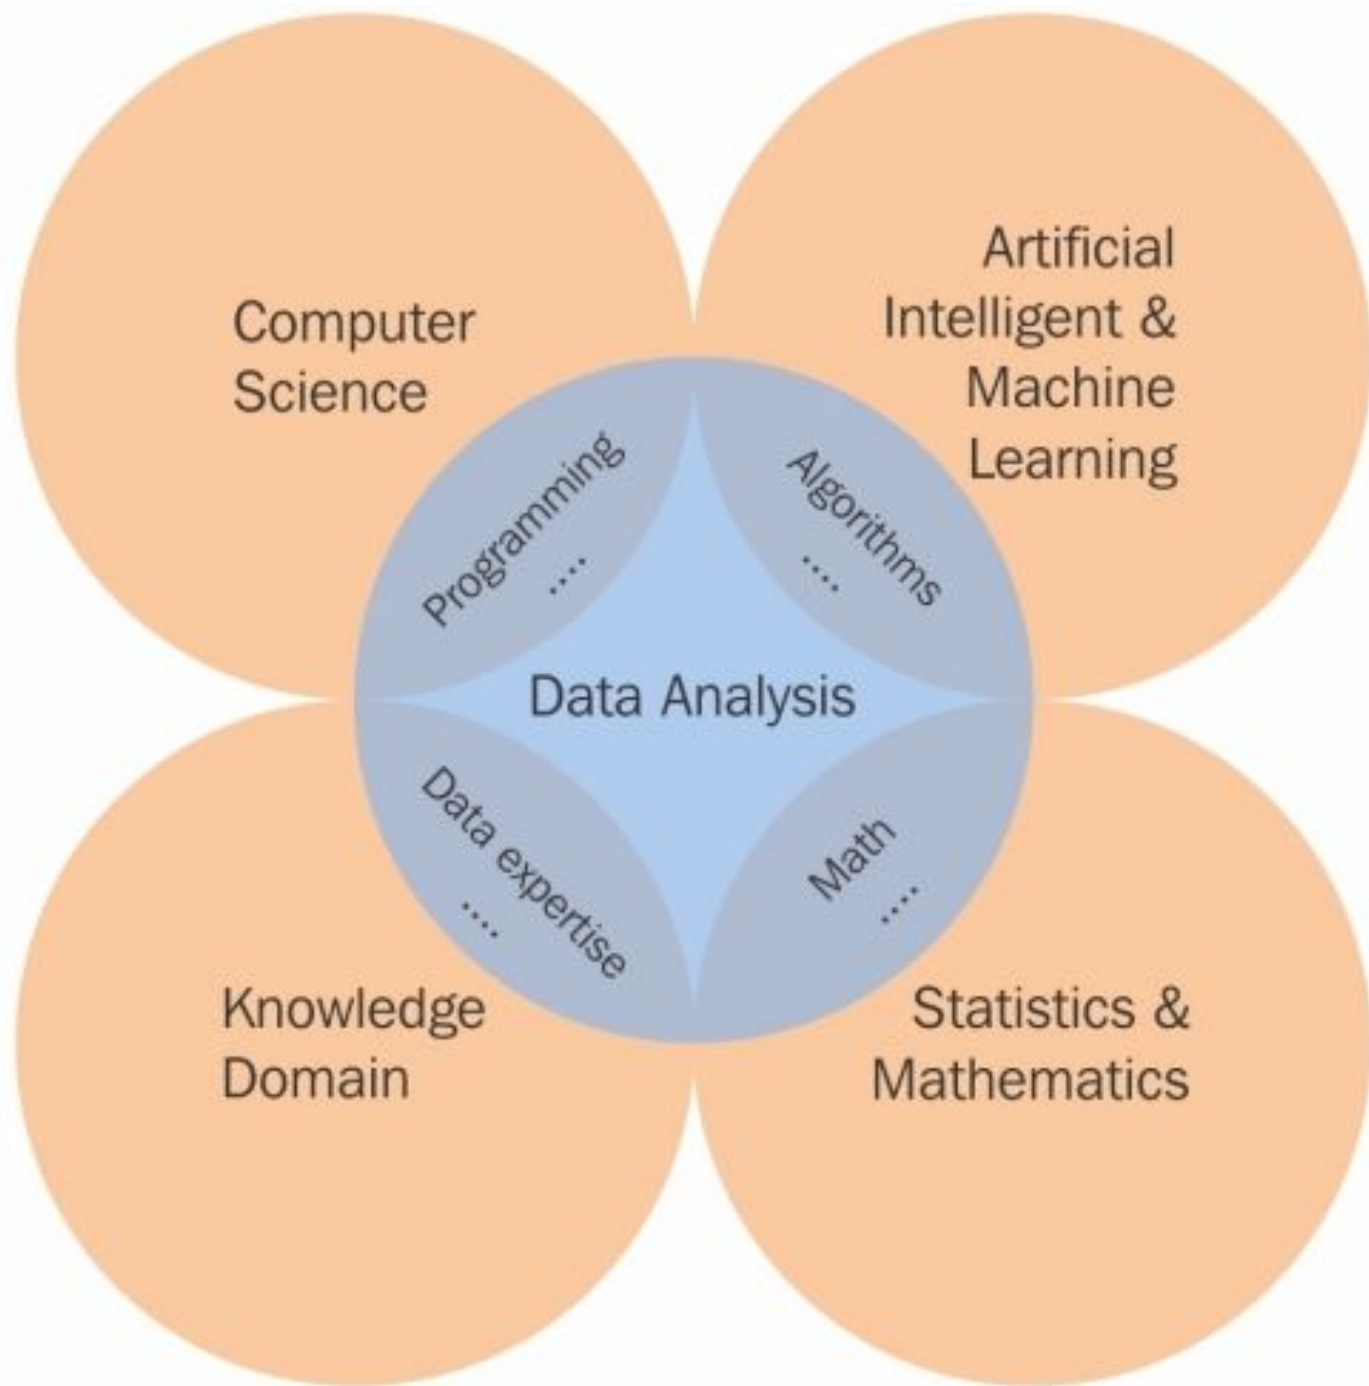
# Google Classroom

- The resources can be found on the Google classroom
- Goto <https://classroom.google.com> and enter the class code below
- If you have problem to access the Google Classroom, please inform trainer or the staff

gdsag5n

# Topic 1

## Data Preparation



# Data Analysis Steps


- Data Preparation / Processing
  - Data Collection
  - Data Pre-Processing
- Data Analytics
  - Data Visualization
  - Data Analysis and Exploration
- Data Modeling
  - Create Model
  - Train (Fit) Model
  - Deploy Model

# Python Libraries for Data Analysis

- Data Processing and Analysis

- Numpy

- Pandas



NumPy




pandas

- Data Visualization

- Matplotlib

- Seaborn

matplotlib



Seaborn

- Data Modeling

- Scikit Learn



# Google Colab for Python

- Google Colaboratory is a free Jupyter notebook environment that requires no setup and runs entirely in the cloud.
- With Colaboratory you can write and execute Python and Tensorflow code, save and share your analyses
- You can also access computing resources such as GPU and TPU for free
- You can access Google Colab by typing

<https://colab.research.google.com>

# Access Google Drive from Colab

To mount google drive locally, refer to this link  
<https://colab.research.google.com/notebooks/io.ipynb>

```
from google.colab import drive  
drive.mount('/content/gdrive')
```

To access your Google Drive :

```
ls '/content/gdrive/My Drive/'
```

# Install Python Packages

`pip install numpy`

`pip install matplotlib`

`pip install pandas`

`pip install scipy`

`pip install sklearn`

For mac user, use `pip3` instead of `pip`



# Import Python Packages

```
import numpy as np  
import matplotlib.pyplot as plt  
import seaborn as sb  
import pandas as pd  
import sklearn
```

# Pandas

- Data processing is important part of analyzing the data, because data is not always available in desired format.
- Various processing are required before analyzing the data such as cleaning, restructuring or merging etc.
- Pandas are built on the top of Numpy.
- Pandas provides rich set of functions to process various types of data
- Pandas integrates well with matplotlib library, which makes it very handy tool for analyzing the data.




# DataFrame

- DataFrame is the key data structure in Python, similar to the Dataframe in R
- DataFrame behave likes Excel spreadsheet
- It has rows index and columns name

# Series

- Pandas provides two very useful data structures to process the data i.e. Series and DataFrame
- Series is a one-dimensional labeled array capable of holding any data type (integers, strings, floating point numbers, Python objects, etc.).
- The axis labels are collectively referred to as the index.
- The basic method to create a Series is to call:  
`s = pd.Series(data, index=index)`
- Example  
`data = np.array([10,20,30,40])`  
`s = pd.Series(data,index=['2011','2012','2013',`



# Create Series

- Series can be instantiated from dicts
- Example

```
data = {'2011':40,'2012':30,'2013':20,'2014':10}  
s = pd.Series(data)
```

# Retrieve Series Data

- A Series is like a fixed-size dict in that you can get and set values by index label.
- Example:


```
data = {'2011':40,'2012':30,'2013':20,'2014':10}
```

```
s = pd.Series(data)
```

```
s['2012']
```

# Data Frame

- DataFrame is the widely used data structure of pandas.
- DataFrame has two different index i.e. column-index and row-index.
- You can think of it as an SQL table or a spreadsheet data representation.



# Create Dataframe


- The most common way to create a DataFrame is by using the dictionary of equal-length list as shown below:

```
data = {  
    'Name' : ["Ally","Belinda","Jane","Steve"],  
    'Height' : [160,165,155,180],  
    'Gender' : ['F','F','F','M']  
}  
df = pd.DataFrame(data)
```



# Import and Export Data

- Pandas supports the integration with many file formats or data sources out of the box (csv, excel, sql, json, parquet,...).
- Importing data from each of these data sources is provided by function with the prefix `read_*`. Similarly, the `to_*` methods are used to store data.



# Import and Export CSV Data

- Pandas provides the `read_csv()` function to read data stored as a csv file into a pandas DataFrame. Eg

```
df = pd.read_csv('mtcars.csv')
```

- To create the index and select data from imported data, Eg

```
df = pd.read_csv('mtcars.csv',  
index_col = 'car_names',  
usecols = ['car_names','mpg','cyl','hp'])
```

- Whereas `read_*` functions are used to read data to pandas, the `to_*` methods are used to store data. Eg `mtcars_sample.to_csv('cars_sample.csv')`

# Import and Export Excel Data

- The `to_excel()` method stores the data as an excel file.
- By setting `index=False` the row index labels are not saved in the spreadsheet. Eg

```
mtcars_sample.to_excel('cars_sample.xlsx',  
sheet_name='cars', index=False)
```

- The equivalent read function `read_excel()` will reload the data to a DataFrame. Eg

```
mtcars_sample2 = pd.read_excel('cars_sample.xlsx',  
sheet_name='cars')
```

# DataFrame Attributes

`mtcars.info()` : Information of the dataframe  
`df.shape` : Shape of a dataframe  
`df.columns` : columns of a dataframe  
`df.index` : index of a a dataframe  
`df['col'].values` : values of a particular column

# Activity: Import Data

- Import the Singapore Health Expenditure dataset from

<https://raw.githubusercontent.com/tertiarycourses/datasets/master/government-health-expenditure.csv>

- Use financial year as index
- Import only operating, development and government health expenditure
- Export the data to csv format

# Head and Tail

- To view a small sample of a Series or the DataFrame object, use the `head()` and the `tail()` methods.
- `head()` returns the first `n` rows (observe the index values). The default number of elements to display is five, but you may pass a custom number.
- `tail()` returns the last `n` rows (observe the index values). The default number of elements to display is five, but you may pass a custom number.
- Example  
`mtcars.head(10)`  
`mtcars.tail()`

# Select Column

- To select a single column, use square brackets `[]` with the column name of the column of interest. The returned data type is a pandas Series Eg `cars_sample['mpg']`
- To select multiple columns, use a list of column names within the selection brackets `[]`. The returned data type is a pandas DataFrame. Eg `cars_sample[['mpg','cyl']]`



# Select Row

- The Python and NumPy indexing operators "[" and attribute operator "." provide quick and easy access to Pandas data structures across a wide range of use cases.
- Pandas now supports three types of Multi-axes indexing; the three types
  - .loc() - Label based
  - .iloc() - Integer based
- Example:
  - mtcars\_sample.loc['Fiat 128']
  - mtcars\_sample.loc[['Fiat 128','Lotus Europa']]
  - mtcars\_sample.iloc[3]
  - mtcars\_sample.iloc[[3,5]]



# Slicing Data

- `iloc` can be used to slice out a subset of the data.

Example

```
mtcars_sample.iloc[3:6]
```

```
mtcars_sample.iloc[:5]
```

# Activity: Selecting and Slicing Data

- Import the Singapore Health Expenditure dataset from


<https://raw.githubusercontent.com/tertiarycourses/datasets/master/government-health-expenditure.csv>

- Retrieve the operating and development expenditure data from the year 2016 and 2017
- Slice out the operating and development expenditure data from 2009 to 2013

# Filtering Data

- To select rows based on a conditional expression, use a condition inside the selection brackets [].
- Some examples are as follows:

```
mtcars_sample[mtcars_sample['cyl']>4]  
mtcars_sample[(mtcars_sample["mpg"] > 20) | (mtcars_sample["cyl"] < 6)]  
mtcars_sample[mtcars_sample["am"] == 1]  
mtcars_sample.loc[["Mazda RX4", "Fiat 128"], :]  
mtcars_sample[mtcars_sample['cyl'].isin([6,8])]
```



# Activity: Filtering Data

- Import the Singapore Health Expenditure dataset from

<https://raw.githubusercontent.com/tertiarycourses/datasets/master/government-health-expenditure.csv>

- Retrieve the all the data where operating expenditure data >5000
- Retrieve the all the data where operating expenditure data is between 5000 and 8000

# Missing Values

- Missing values are represented with NaN in Pandas, for example,

	<b>one</b>	<b>two</b>	<b>three</b>
<b>a</b>	1.102077	2.012164	0.072745
<b>b</b>	NaN	NaN	NaN
<b>c</b>	-0.011272	0.361001	-0.821974
<b>d</b>	NaN	NaN	NaN
<b>e</b>	-0.090309	0.553269	-0.065935
<b>f</b>	0.792010	0.028055	0.524832
<b>g</b>	NaN	NaN	NaN
<b>h</b>	-0.036673	-2.037336	-0.595914

# Remove Missing Data

- Use `isnull()` to check any missing data  
`df['one'].isnull()`
- Use `dropna()` to remove the missing  
`df.dropna()`

# Impute Missing Data

- You can impute missing data with a fixed number, or forward fill or backfill

```
df.fillna(0)
```

```
df.fillna(method='pad')
```

```
df.fillna(method='backfill')
```

# Activity: Filtering Data

- Import the Singapore Hospital Admission dataset from

<https://raw.githubusercontent.com/tertiarycourses/datasets/master/hospital-admissions-by-sector-annual.csv>

- Remove the missing data with 'na'

	year	level_1	level_2	value
0	1984	Acute Hospitals Admissions	Public	na
1	1984	Acute Hospitals Admissions	Non-public	na
2	1984	Psychiatric Hospitals Admissions	Public	na
3	1984	Psychiatric Hospitals Admissions	Non-public	na
4	1984	Community Hospitals Admissions	Public	na
...	...	...	...	...
211	2019	Acute Hospitals Admissions	Non-public	134197
212	2019	Psychiatric Hospitals Admissions	Public	9234
213	2019	Psychiatric Hospitals Admissions	Non-public	0
214	2019	Community Hospitals Admissions	Public	10215
215	2019	Community Hospitals Admissions	Non-public	9828

216 rows × 4 columns






# Topic 2

## Data Transformation

# Concatenating Data

- The `concat()` function performs concatenation operations of multiple tables along one of the axis (0:row-wise or 1:column-wise).
- By default concatenation is along axis 0, so the resulting table combines the rows of the input tables.
- The syntax is  
`pd.concat(objs, axis=0, join='outer', ignore_index=False, keys=None, levels=None, names=None, verify_integrity=False, copy=True)`



# Concatenating Data Demo

```
merc = [c for c in mtcars_sample.index if 'Merc' in c]
merc_cars = mtcars_sample.loc[merc]
toyota = [c for c in mtcars_sample.index if 'Toyota' in c]
toyota_cars = mtcars_sample.loc[toyota]
merc_toyota_cars = pd.concat([merc_cars,
toyota_cars], axis=0)
```

merc_cars				toyota_cars			
car_names	mpg	cyl	hp	car_names	mpg	cyl	hp
Merc 240D	24.4	4	62	Merc 240D	24.4	4	62
Merc 230	22.8	4	95	Merc 230	22.8	4	95
Merc 280	19.2	6	123	Merc 280	19.2	6	123
Merc 280C	17.8	6	123	Merc 280C	17.8	6	123
Merc 450SE	16.4	8	180	Merc 450SE	16.4	8	180
Merc 450SL	17.3	8	180	Merc 450SL	17.3	8	180
Merc 450SLC	15.2	8	180	Merc 450SLC	15.2	8	180
				Toyota Corolla	33.9	4	65
				Toyota Corona	21.5	4	97

# Activity: Selecting and Slicing Data

- Import the air quality and pm2.5 data as follows:  

```
air_quality_no2 =  
pd.read_csv("https://raw.githubusercontent.com/pandas-dev/pandas/master/doc/data/air_quality_no2_long.csv")  
air_quality_no2 = air_quality_no2[["date.utc", "location", "parameter", "value"]]  
  
air_quality_pm25 =  
pd.read_csv("https://raw.githubusercontent.com/pandas-dev/pandas/master/doc/data/air_quality_pm25_long.csv")  
air_quality_pm25 = air_quality_pm25[["date.utc", "location", "parameter", "value"]]
```
- Join the two datasets row wise.

# Appending Data

Alternate way to join the data is to use the append function. The syntax is as follows:


```
toyota_merc_cars2= toyota_cars.append(merc_cars)
toyota_merc_cars2
```

merc_cars				toyota_cars				toyota_merc_cars2			
car_names	mpg	cyl	hp	car_names	mpg	cyl	hp	car_names	mpg	cyl	hp
Merc 240D	24.4	4	62					Merc 240D	24.4	4	62
Merc 230	22.8	4	95					Merc 230	22.8	4	95
Merc 280	19.2	6	123					Merc 280	19.2	6	123
Merc 280C	17.8	6	123					Merc 280C	17.8	6	123
Merc 450SE	16.4	8	180					Merc 450SE	16.4	8	180
Merc 450SL	17.3	8	180					Merc 450SL	17.3	8	180
Merc 450SLC	15.2	8	180					Merc 450SLC	15.2	8	180
				Toyota Corolla	33.9	4	65	Toyota Corolla	33.9	4	65
				Toyota Corona	21.5	4	97	Toyota Corona	21.5	4	97


# Merging Data

- Pandas has full-featured, high performance in-memory join operations idiomatically very similar to relational databases like SQL.
- The syntax is  
`pd.merge(left, right, how='inner', on=None, left_on=None, right_on=None, left_index=False, right_index=False, sort=True, suffixes=('_x', '_y'), copy=True, indicator=False, validate=None)`


INNER JOIN




FULL JOIN



LEFT JOIN



RIGHT JOIN



# Merging Data Demo

```
left = pd.DataFrame({'key': ['K0', 'K1', 'K2', 'K3'],  
                    'A': ['A0', 'A1', 'A2', 'A3'],  
                    'B': ['B0', 'B1', 'B2', 'B3']})  
right = pd.DataFrame({'key': ['K0', 'K1', 'K2', 'K4'],  
                     'C': ['C0', 'C1', 'C2', 'C4'],  
                     'D': ['D0', 'D1', 'D2', 'D4']})  
result = pd.merge(left, right, on='key', how='inner')
```

key    A    B					key    C    D					key    A    B    C    D					
0	K0	A0	B0	+	0	K0	C0	D0	=	0	K0	A0	B0	C0	D0
1	K1	A1	B1		1	K1	C1	D1		1	K1	A1	B1	C1	D1
2	K2	A2	B2		2	K2	C2	D2		2	K2	A2	B2	C2	D2
3	K3	A3	B3		3	K4	C4	D4							

# Activity: Merging Data

Merge air quality and pm 2.5 data using inner join

- based on location
- based on date




# Sorting Data

- With `sort_values()`, the rows in the table are sorted according to the defined column(s). The index will follow the row order.
- Example  
`mtcars_sample.sort_values(by="cyl",ascending=False)`

# Group and Aggregate Data

Index	Fruit
1	Apple
2	Apple
3	Orange
4	Apple
5	Orange



Fruit	
Apple	3
Orange	2

- To compare subsets
- To deduce reasons why subgroups differ
- To subset your data for your analysis

# Groupby

- A groupby operation involves some combination of splitting the object, applying a function, and combining the results.

```
mtcars_sample.groupby(['cyl']).mean()
```

```
mtcars_sample.groupby('cyl').hp.mean()
```

```
mtcars_sample.groupby(['cyl']).sum()
```

```
mtcars_sample.groupby(['cyl']).agg(['mean', 'count'])
```

```
mtcars_sample.groupby(['cyl','am']).mean()
```

```
mtcars_sample.groupby('cyl').agg(lambda x:max(x)-min(x))
```

```
mtcars_sample.groupby(['cyl', 'am']).agg(['mean', 'count'])
```

						mpg	hp
				cyl	am		
cyl				4	0	22.900000	84.666667
					1	28.075000	81.875000
				6	0	19.125000	115.250000
					1	20.566667	131.666667
				8	0	15.050000	194.166667
					1	15.400000	299.500000

				mpg	hp	am
cyl						
4	293.3	909	8			
6	138.2	856	3			
8	211.4	2929	2			

# Activity: Groupby

- Import the Singapore long term care facilities data  
<https://raw.githubusercontent.com/tertiarycourses/datasets/master/number-of-residential-long-term-care-facilities-sector-breakdown.csv>
- Compute the total number of long term care facilities breakdown by year and sector using Group By method

Time: 10 mins

# Pivot Table

- Pivot table is a well known concept in spreadsheet to reshape the data.
- `pivot()` can be used to rearrange the data
- `pivot_table()` can be used, providing an aggregation function (e.g. mean) on how to combine these values.

df

	foo	bar	baz	zoo
0	one	A	1	x
1	one	B	2	y
2	one	C	3	z
3	two	A	4	q
4	two	B	5	w
5	two	C	6	t



```
df.pivot(index='foo',  
          columns='bar',  
          values='baz')
```

bar	A	B	C
foo			
one	1	2	3
two	4	5	6

# Pivot Table

```
mtcars_sample.pivot(columns='cyl',values='hp')  
mtcars_sample.pivot(columns='cyl',values='hp').mean()
```

```
   cyl  
4    82.636364  
6   122.285714  
8   209.214286  
dtype: float64
```

```
mtcars_sample.pivot_table(index='cyl',columns='am',  
values='hp',aggfunc='mean')
```

	am	
	0	1
cyl		
4	84.666667	81.875000
6	115.250000	131.666667
8	194.166667	299.500000

# Activity: Pivot Table

- Import the Singapore long term care facilities data  
<https://raw.githubusercontent.com/tertiarycourses/datasets/master/number-of-residential-long-term-care-facilities-sector-breakdown.csv>
- Compute the total number of long term care facilities breakdown by year and sector using Pivot Table method

Time: 10 mins


# Topic 3

## Data Visualization




# Create Plots in Pandas

- You can quickly plot the dataframe data using the plot method
- By default, plot will yield a line plot
- Example:  
`mtcars_sample.plot()`




`.plot.*`



# Activity: Line Plot

Plot the Singapore healthcare expenditure over the year




# Pandas Plots

Pandas support the following data visualizations:

- Area plot: `area`
- Bar plot: `bar`
- Horizontal Bar plot: `barh`
- Boxplot: `box`
- Density plot: `density`
- Histogram: `hist`
- Line plot: `line`
- Pie plot: `pie`
- Scatter plot: `scatter`

# Scatter Plot

```
mtcars_sample.plot.scatter(x="hp", y="mpg", alpha=0.5)
```





# Bar Plot

```
mtcars_sample.pivot(columns='cyl',values='hp').mean().plot.bar()
```


```
mtcars_sample.pivot(columns='cyl',values='hp').mean().plot.bar(color='red')
```

```
mtcars_cyl=  
mtcars_sample.pivot(columns='cyl',values='hp').mean().plot(kind='bar',color='red')
```




# Horizontal Bar Plot

```
mtcars_cyl=  
mtcars_sample.pivot(columns='cyl',values='hp').mean().plot.barh()
```




# Stacked Bar Plot

```
mtcars_sample.pivot_table(index='am',columns='cyl',values='mpg',aggfunc=['mean']).plot.bar()
```



# Stacked Bar Plot


```
mtcars_sample.groupby(['am','cyl']).mean().plot.bar(stacked=True)
```






# Pie Plot

```
mtcars_cyl=  
mtcars_sample.pivot(columns='cyl',values='hp').mean().plot.pie()
```



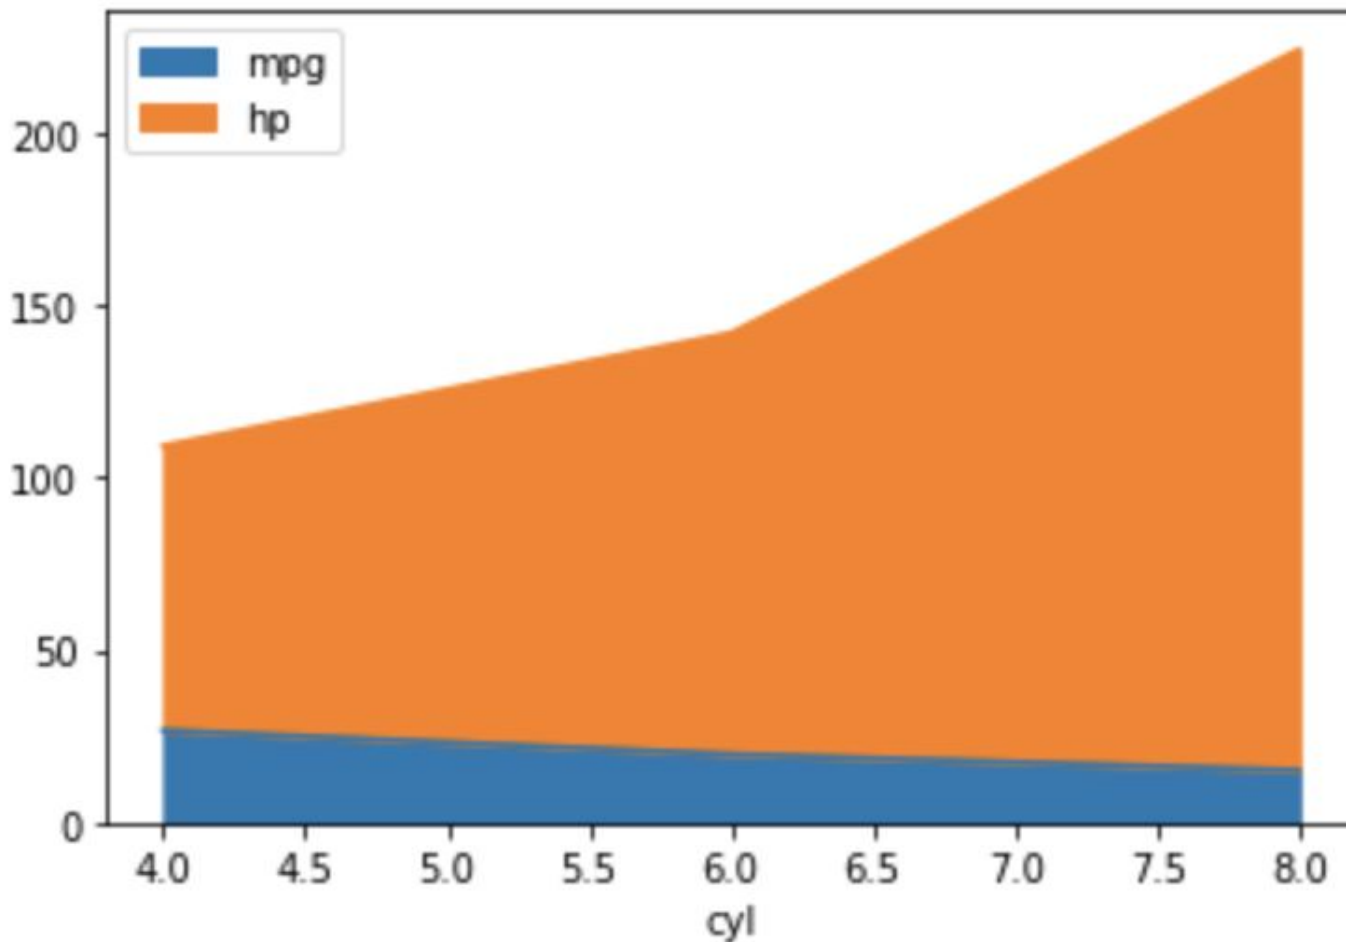
# Box Plot

```
mtcars_cyl= mtcars_sample.pivot(columns='cyl',values='hp').plot.box()
```




# Area Plot

```
mtcars_sample2 = mtcars_sample[['cyl','mpg','hp']]  
mtcars_sample2.groupby(['cyl']).mean().plot.area()
```



# Histogram

```
mtcars_sample.mpg.plot.hist(bins=15)
```




# Activity: Data Visualization

- Import the Singapore long term care facilities data  
<https://raw.githubusercontent.com/tertiarycourses/datasets/master/number-of-residential-long-term-care-facilities-sector-breakdown.csv>
- Create a horizontal bar plot of the total long term care facilities by the sector

Time: 10 mins

# Subplot

```
mtcars_cyl= mtcars_sample.groupby('cyl').mean()  
mtcars_cyl.plot(subplots=True)
```




# Activity: Data Visualization

- Import the Singapore Health Expenditure dataset from

<https://raw.githubusercontent.com/tertiarycourses/datasets/master/government-health-expenditure.csv>


- Create 3 subplots for each expenditure.



# Matplotlib

- Panda Series and DataFrame data structures work seamlessly with more advanced data visualization tools such as Matplotlib or Seaborn.
- For example, you can do a scatter plot using Matplotlib as follows

```
plt.scatter(x="hp", y="mpg", data = mtcars_sample,color='red')  
plt.xlabel('hp')  
plt.ylabel('mpg')  
plt.title('hp vs mpg')
```







# Seaborn

- You can create heatmap and regression plots readily with Seaborn.

```
mtcars_cyl= mtcars_sample.groupby('cyl').mean()  
sb.heatmap(mtcars_cyl, annot=True)  
sb.lmplot(x="hp", y="mpg",data=mtcars_sample,fit_reg=True)
```



# Topic 4

## Data Analysis

# What is Statistics?

Statistics is a discipline which is concerned with:

- designing experiments and other data collection,
- summarizing information to aid understanding,
- drawing conclusions from data, and
- estimating the present or predicting the future.

Statistical statements:

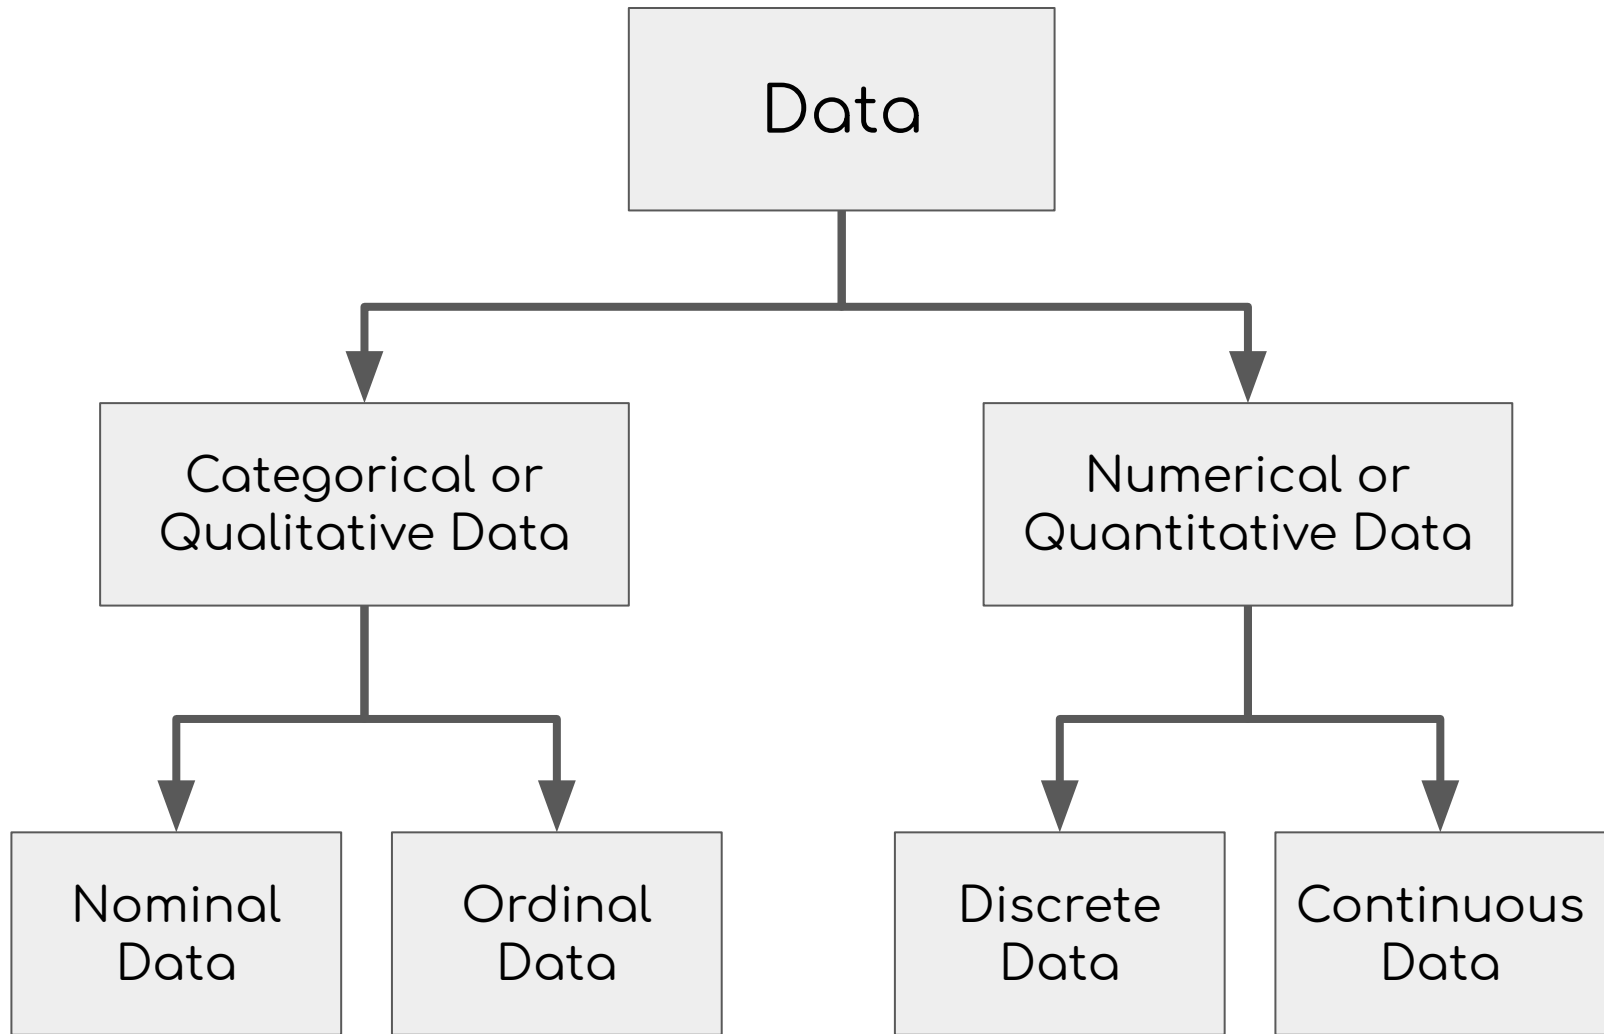
"I sleep for about eight hours per night on average"

"You are more likely to pass the exam if you start preparing earlier"

# Why Statistics Matter?

- Environmental Study
  - Is Singapore getting hotter over last 10 years?
- Policy Study
  - Is more people using green transport such as Bicycles, Buses, Carpool, CNG Cars, Electric Cars, Electric Scooters
- Market Analysis
  - Is more people likely to take green transport if they've seen a recent TV advertisement for green transport?
- Public Transport
  - Is more people likely to commute by MRT if we have more MRT stations in the neighborhood?
- Health Care
  - Does air pollution from vehicles cause any health concern?
- Data Science
  - Statistics is fundamental for understanding Artificial Intelligence and Machine Learning.

# Types of Data



# Categorical and Quantitative Data

- **Categorical (Qualitative) Data** - each observation belongs to one of a set of categories. Examples:
  - Weather (Rainy /Sunny)
  - Air Pollutants (Ozone/Nitrogen Dioxide)
  - Gender (Male or Female)
  - Place of residence (HDB, Condo, ...)
  - Marital status (Married, Single,...)
- **Quantitative (Numerical) Data** - observations take numerical values. Examples:
  - Surface Air temperature
  - Weekly number of dengue cases
  - No. of days with rainfall in a month
  - Age
  - Number of cars
  - Weight

# Nominal and Ordinal Data

- **Nominal Data** is defined as data that is used for naming or labelling variables, without any quantitative value. It is sometimes called “labels” data Eg
  - Male/Female
  - Red/Green/Blue
- **Ordinal Data** is a type of categorical data with an order. The variables in ordinal data are listed in an ordered manner.
  - Disagree/Neutral/Agree/Strongly Agree
  - Very Bad/Bad/Good/Very Good

# Discrete and Continuous Data


- Discrete Data is a set of countable numbers such as 0, 1, 2, 3,.....Examples:
  - No. of days with rainfall in a month
  - Weekly no. of dengue cases
  - Number of children in a family
  - Number of foreign languages spoken
- Continuous Data are continuous numbers from an interval. Examples:
  - Surface Air temperature
  - Amount of rainfall in a month
  - Height
  - Weight



# Measures of Central Tendency


## Central Tendency

### Mean




The center of gravity or the balance point

### Median



Midpoint of ranked values

### Mode




Most frequently observed value

- Mean - add up all the values and divide by how many there are
- Median - Arrange all the numbers from smallest to largest:
  - odd number of points: Median = middle value
  - even number of points: Median = mean of the middle two values


# Mean vs Median

- Mean
  - Useful for roughly symmetric quantitative data
  - Sensitive to outlier data
- Median
  - Splits the data into halves
  - Useful for highly skewed quantitative data
  - Insensitive to outlier data



# Measures of Dispersion

- The measures of dispersion measure the differences between how far “spread out” the data values are.
- Two commonly used measures for dispersion are: range and standard deviation.
- 



**Same center,  
different variation**


# Standard Deviation

- The standard deviation measures the dispersion of a dataset relative to its mean and is calculated as the square root of the variance
- Larger standard deviation means greater variability of the data.


$$\text{SD} = \sqrt{\frac{\sum |x - \bar{x}|^2}{n}}$$

# Range

- Range is the difference between the highest and lowest values.
- Since it uses only the extreme values, it is greatly affected by extreme values.
- Range ignores the way in which data are distributed



$$\text{Range} = 12 - 7 = 5$$



$$\text{Range} = 12 - 7 = 5$$


# Statistical Functions in Pandas

Pandas offer the following built in statistical functions:

- `count()`      Number of non-null observations
- `sum()`      Sum of values
- `mean()`      Mean of Values
- `median()`      Median of Values
- `mode()`      Mode of values
- `std()`      Standard Deviation of the Values
- `min()`      Minimum Value
- `max()`      Maximum Value
- `abs()`      Absolute Value
- `prod()`      Product of Values
- `cumsum()`      Cumulative Sum
- `cumprod()`      Cumulative Product

# Aggregating Statistics

- Different statistics are available and can be applied to columns with numerical data.
- Aggregating statistical operations in general exclude missing data and operate across rows by default
- For example,  
`mtcars_sample[["mpg", "hp"]].median()`



# Descriptive Statistics

- You can get the descriptive statistics of the data using the describe function
- For example,

```
mtcars_sample[["mpg", "hp"]].median()  
mtcars_sample[["mpg", "hp"]].describe()
```

	<b>mpg</b>	<b>hp</b>
<b>count</b>	32.000000	32.000000
<b>mean</b>	20.090625	146.687500
<b>std</b>	6.026948	68.562868
<b>min</b>	10.400000	52.000000
<b>25%</b>	15.425000	96.500000
<b>50%</b>	19.200000	123.000000
<b>75%</b>	22.800000	180.000000
<b>max</b>	33.900000	335.000000





# Count

- You can use `value_count()` to count the number of records in each category
- For example  
`mtcars_sample["cyl"].value_counts()`

```
8      14
```

```
4      11
```

```
6       7
```

```
Name: cyl, dtype: int64
```

# Activity: Descriptive Statistics

- Import the Singapore long term care facilities data  
<https://raw.githubusercontent.com/tertiarycourses/datasets/master/number-of-residential-long-term-care-facilities-sector-breakdown.csv>
- Compute the number of facilities in each sector using the `value_counts()` method

Time: 10 mins

# What is Covariance

- Variance is a measure of the variability or spread in a set of data
- We use the following formula to compute variance for population and sample respectively.

$$Var(x) = \frac{\sum (x - \bar{x})^2}{N} \quad Var(x) = \frac{\sum (x - \bar{x})^2}{N - 1}$$

- **Covariance** is a measure of the extent to which corresponding elements from two sets of ordered data move in the same direction.
- We use the following formula to compute covariance for population and sample respectively





$$Cov(x, y) = \frac{\sum (x - \bar{x})(y - \bar{y})}{N} \quad Cov(x, y) = \frac{\sum (x - \bar{x})(y - \bar{y})}{N - 1}$$

# Covariance Matrix

- Variance and covariance are often displayed together in a covariance matrix given as follows:

$$\text{Cov}(A) = \begin{bmatrix} \frac{\sum (x_i - \bar{X})(x_i - \bar{X})}{N} & \frac{\sum (x_i - \bar{X})(y_i - \bar{Y})}{N} \\ \frac{\sum (x_i - \bar{X})(y_i - \bar{Y})}{N} & \frac{\sum (y_i - \bar{Y})(y_i - \bar{Y})}{N} \end{bmatrix}$$
$$= \begin{bmatrix} \text{Cov}(X, X) & \text{Cov}(Y, X) \\ \text{Cov}(X, Y) & \text{Cov}(Y, Y) \end{bmatrix}$$

# Covariance Matrix Visualization



# Covariance on Pandas

- You can use cov function to compute the covariance matrix
- For example,  
`mtcars_sample.cov()`

	<b>mpg</b>	<b>cyl</b>	<b>hp</b>	<b>am</b>
<b>mpg</b>	36.324103	-9.172379	-320.732056	1.803931
<b>cyl</b>	-9.172379	3.189516	101.931452	-0.465726
<b>hp</b>	-320.732056	101.931452	4700.866935	-8.320565
<b>am</b>	1.803931	-0.465726	-8.320565	0.248992

# What is Correlation

- The correlation coefficient is also known as the Pearson product-moment correlation coefficient, or Pearson's correlation coefficient.
- It is obtained by dividing the covariance of the two variables by the product of their standard deviations.

$$\text{Corr}(x, y) = \frac{\text{Cov}(x, y)}{\sigma_x \sigma_y}$$

- The values of the correlation coefficient can range from -1 to +1. The closer it is to +1 or -1, the more closely are the two variables are related.
- The positive sign signifies the direction of the correlation i.e. if one of the variables increases, the other variable is also supposed to increase.







# Correlation Matrix

- For multiple variables, we can display all the correlation coefficients in the matrix form as below:

$$\begin{bmatrix} 1 & \text{Corr}(X,Y) & \text{Corr}(X,Z) \\ \text{Corr}(X,Y) & 1 & \text{Corr}(Y,Z) \\ \text{Corr}(X,Z) & \text{Corr}(Y,Z) & 1 \end{bmatrix}$$

# Correlation Coefficient



# Correlation Matrix on Pandas

- You can use cov function to compute the covariance matrix
- For example,  
`mtcars_sample.cov()`

	<b>mpg</b>	<b>cyl</b>	<b>hp</b>	<b>am</b>
<b>mpg</b>	36.324103	-9.172379	-320.732056	1.803931
<b>cyl</b>	-9.172379	3.189516	101.931452	-0.465726
<b>hp</b>	-320.732056	101.931452	4700.866935	-8.320565
<b>am</b>	1.803931	-0.465726	-8.320565	0.248992

# Pandas Datetime

- By applying the `to_datetime` function, pandas interprets the strings and convert these to datetime (i.e. `datetime64[ns, UTC]`) objects.
- In pandas we call these datetime objects similar to `datetime.datetime` from the standard library as `pandas.Timestamp`

```
air_quality =  
pd.read_csv("https://raw.githubusercontent.com/pandas-dev/pandas/master/doc/data/air_quality_no2_long.csv", parse_dates=["date.utc"])
```

- Using `pandas.Timestamp` for datetimes enables us to calculate with date information and make them comparable. Hence, we can use this to get the length of our time series Eg

```
air_quality["datetime"].max() - air_quality["datetime"].min()
```

# Datetime Properties


- By using Timestamp objects for dates, a lot of time-related properties are provided by pandas. For example the month, but also year, weekofyear, quarter
- All of these properties are accessible by the dt accessor. For example

```
air_quality["month"] = air_quality["datetime"].dt.month  
air_quality.groupby([air_quality["datetime"].dt.weekday,"location"])  
["value"].mean()
```

```
datetime  location  value  
0         BETR801    27.875000  
          FR04014    24.856250  
          London Westminster  23.969697  
1         BETR801    22.214286  
          FR04014    30.999359  
          London Westminster  24.885714  
2         BETR801    21.125000  
          FR04014    29.165753  
          London Westminster  23.460432  
3         BETR801    27.500000  
          FR04014    28.600690  
          London Westminster  24.780142  
4         BETR801    28.400000  
          FR04014    31.617986  
          London Westminster  26.446809  
5         BETR801    33.500000  
          FR04014    25.266154  
          London Westminster  24.977612  
6         BETR801    21.896552  
          FR04014    23.274306  
          London Westminster  24.859155  
Name: value, dtype: float64
```

# Time Series Plot

- We can calculate a given statistic (e.g. mean) for each hour of the day
- We can use the groupby and datetime property hour of pandas Timestamp, which is also accessible by the dt accessor.




# Datetime as Index

- Working with a datetime index (i.e. DatetimeIndex) provides powerful functionalities.
- For example, we do not need the dt accessor to get the time series properties, but have these properties available on the index directly

```
no_2 = air_quality.pivot(index="datetime", columns="location",  
values="value")
```

```
no_2["2019-05-20":"2019-05-21"].plot()
```




# Resample a Time Series

- A very powerful method on time series data with a datetime index, is the ability to `resample()` time series to another frequency (e.g., converting secondly data into 5-minutely data).
- The `resample()` method is similar to a `groupby` operation:
  - it provides a time-based grouping, by using a string (e.g. M, 5H,...) that defines the target frequency
  - it requires an aggregation function such as `mean`, `max`,...




# Resampled Time Series Plot

```
no_2.resample("D").mean().plot(style="-o", figsize=(10, 5));
```



# Activity: Time Series Analysis

- Import the Singapore long term care facilities data <https://raw.githubusercontent.com/tertiarycourses/datasets/master/number-of-residential-long-term-care-facilities-sector-breakdown.csv>
- Plot the total no of long term care facilities vs year



# Summary Q&A



**Practice  
Makes  
Perfect**

# Feedback

<https://goo.gl/R2eumq>



# Thank You!

Marcus Lee Yi Qing

87119800

[makasulee@gmail.com](mailto:makasulee@gmail.com)